

In [2]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import datetime as dt
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
import seaborn as sns
# Load dataset
df = pd.read_csv('C:/Users/Adamin/OneDrive/Desktop/sales_data_sample.csv', encoding='utf-8')
df.head()
```

Out[2]:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	OR
0	10107	30	95.70	2	2871.00	1
1	10121	34	81.35	5	2765.90	2
2	10134	41	94.74	2	3884.34	3
3	10145	45	83.26	6	3746.70	4
4	10159	49	100.00	14	5205.27	10

5 rows × 25 columns



In [3]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   ORDERNUMBER      2823 non-null   int64  
 1   QUANTITYORDERED 2823 non-null   int64  
 2   PRICEEACH        2823 non-null   float64 
 3   ORDERLINENUMBER 2823 non-null   int64  
 4   SALES            2823 non-null   float64 
 5   ORDERDATE        2823 non-null   object  
 6   STATUS            2823 non-null   object  
 7   QTR_ID           2823 non-null   int64  
 8   MONTH_ID         2823 non-null   int64  
 9   YEAR_ID          2823 non-null   int64  
 10  PRODUCTLINE      2823 non-null   object  
 11  MSRP              2823 non-null   int64  
 12  PRODUCTCODE      2823 non-null   object  
 13  CUSTOMERNAME     2823 non-null   object  
 14  PHONE             2823 non-null   object  
 15  ADDRESSLINE1     2823 non-null   object  
 16  ADDRESSLINE2     302 non-null    object  
 17  CITY              2823 non-null   object  
 18  STATE             1337 non-null   object  
 19  POSTALCODE        2747 non-null   object  
 20  COUNTRY           2823 non-null   object  
 21  TERRITORY         1749 non-null   object  
 22  CONTACTLASTNAME  2823 non-null   object  
 23  CONTACTFIRSTNAME 2823 non-null   object  
 24  DEALSIZE          2823 non-null   object  
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```

```
In [4]: # Drop unnecessary columns
to_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATE', 'POSTALCODE', 'PHONE']
df = df.drop(to_drop, axis=1)
#Check for null values
df.isnull().sum()
```

```
Out[4]: ORDERNUMBER      0
QUANTITYORDERED      0
PRICEEACH            0
ORDERLINENUMBER      0
SALES                0
ORDERDATE            0
STATUS               0
QTR_ID               0
MONTH_ID             0
YEAR_ID              0
PRODUCTLINE          0
MSRP                 0
PRODUCTCODE          0
CUSTOMERNAME         0
CITY                 0
COUNTRY              0
TERRITORY            1074
CONTACTLASTNAME     0
CONTACTFIRSTNAME    0
DEALSIZE             0
dtype: int64
```

```
In [5]: df.dtypes
```

```
Out[5]: ORDERNUMBER      int64
QUANTITYORDERED    int64
PRICEEACH        float64
ORDERLINENUMBER    int64
SALES            float64
ORDERDATE        object
STATUS            object
QTR_ID           int64
MONTH_ID          int64
YEAR_ID           int64
PRODUCTLINE      object
MSRP              int64
PRODUCTCODE      object
CUSTOMERNAME     object
CITY              object
COUNTRY           object
TERRITORY         object
CONTACTLASTNAME   object
CONTACTFIRSTNAME  object
DEALSIZE          object
dtype: object
```


```
In [9]: # Select numeric columns only
df_numeric = df.select_dtypes(include=['int64', 'float64'])

# Visualize outliers
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(12, 6))
sns.boxplot(data=df_numeric)
plt.title("Outlier Detection using Boxplots (Numeric Columns Only)")
plt.show()

# Identify outliers using IQR
Q1 = df_numeric.quantile(0.25)
Q3 = df_numeric.quantile(0.75)
IQR = Q3 - Q1

# Count outliers per column
outliers = ((df_numeric < (Q1 - 1.5 * IQR)) | (df_numeric > (Q3 + 1.5 * IQR))).sum()
print("\nNumber of Outliers per Numeric Feature:\n", outliers)
```



Number of Outliers per Numeric Feature:

```

ORDERNUMBER      0
QUANTITYORDERED 8
PRICEEACH        0
ORDERLINENUMBER  0
SALES           81
QTR_ID          0
MONTH_ID         0
YEAR_ID          0
MSRP            28
dtype: int64

```

```
In [13]: #normalization data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(df_numeric)
print(" Data normalized using StandardScaler.")

df_normalized = pd.DataFrame(X_scaled, columns=df_numeric.columns)

print("\nSample of Normalized Data:")
display(df_normalized.head())

print("\nMean of each feature after normalization:\n", df_normalized.mean())
print("\nStandard deviation of each feature after normalization:\n", df_normalized.std())
```

Data normalized using StandardScaler.

Sample of Normalized Data:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	C
0	-1.647947	-0.522891	0.596978	-1.057059	-0.370825	-1.4
1	-1.495888	-0.112201	-0.114450	-0.347015	-0.427897	-0.!
2	-1.354689	0.606505	0.549384	-1.057059	0.179443	0.2
3	-1.235214	1.017195	-0.019759	-0.110334	0.104701	0.2
4	-1.083154	1.427884	0.810158	1.783116	0.896740	1.0

Mean of each feature after normalization:

```
ORDERNUMBER      -5.927482e-15
QUANTITYORDERED 3.347580e-16
PRICEEACH        -3.221731e-16
ORDERLINENUMBER  9.312817e-17
SALES            1.812224e-16
QTR_ID           7.550932e-17
MONTH_ID          -1.761884e-17
YEAR_ID           -1.988412e-15
MSRP              8.054328e-17
```

dtype: float64

Standard deviation of each feature after normalization:

```
ORDERNUMBER      1.000177
QUANTITYORDERED 1.000177
PRICEEACH        1.000177
ORDERLINENUMBER  1.000177
SALES            1.000177
QTR_ID           1.000177
MONTH_ID          1.000177
YEAR_ID           1.000177
MSRP              1.000177
```


dtype: float64

```
In [14]: inertia = []
K = range(1, 11)

for k in K:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(X_scaled)
    inertia.append(kmeans.inertia_)

# Plot the Elbow curve
plt.figure(figsize=(8, 5))
plt.plot(K, inertia, marker='o')
plt.title("Elbow Method to Determine Optimal K")
plt.xlabel("Number of Clusters (K)")
plt.ylabel("Inertia")
plt.grid(True)
plt.show()
```

```
C:\Users\Adamin\AppData\Roaming\Python\Python313\site-packages\joblib\externals\loky\backend\context.py:136: UserWarning: Could not find the number of physical cores for the following reason:  
[WinError 2] The system cannot find the file specified  
Returning the number of logical cores instead. You can silence this warning by setting LOKY_MAX_CPU_COUNT to the number of cores you want to use.  
    warnings.warn(  
    File "C:\Users\Adamin\AppData\Roaming\Python\Python313\site-packages\joblib\externals\loky\backend\context.py", line 257, in _count_physical_cores  
        cpu_info = subprocess.run(  
            "wmic CPU Get NumberOfCores /Format:csv".split(),  
            capture_output=True,  
            text=True,  
        )  
    File "C:\Program Files\Python313\Lib\subprocess.py", line 554, in run  
        with Popen(*popenargs, **kwargs) as process:  
            ~~~~~^~~~~~  
    File "C:\Program Files\Python313\Lib\subprocess.py", line 1039, in __init__  
        self._execute_child(args, executable, preexec_fn, close_fds,  
~~~~~^~~~~~^~~~~~^~~~~~^~~~~~^~~~~~^~~~~~  
            pass_fds, cwd, env,  
~~~~~^~~~~~^~~~~~^~~~~~  
            start_new_session, process_group)  
~~~~~^~~~~~^~~~~~^~~~~~^~~~~~  
    File "C:\Program Files\Python313\Lib\subprocess.py", line 1554, in _execute_child  
        hp, ht, pid, tid = _winapi.CreateProcess(executable, args,  
~~~~~^~~~~~^~~~~~^~~~~~^~~~~~  
            # no special security  
~~~~~^~~~~~^~~~~~^~~~~~  
        ...<4 lines>...  
            cwd,  
~~~~~  
            startupinfo)  
~~~~~^~~~~~
```




```
In [15]: from sklearn.metrics import silhouette_score

for k in range (2, 11):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    labels = kmeans.fit_predict(X_scaled)
    sil_score = silhouette_score(X_scaled, labels)
    print(f"K = {k} → Silhouette Score = {sil_score:.4f}")
```

```
K = 2 → Silhouette Score = 0.2169
K = 3 → Silhouette Score = 0.1960
K = 4 → Silhouette Score = 0.2059
K = 5 → Silhouette Score = 0.1872
K = 6 → Silhouette Score = 0.1949
K = 7 → Silhouette Score = 0.1922
K = 8 → Silhouette Score = 0.1837
K = 9 → Silhouette Score = 0.1832
K = 10 → Silhouette Score = 0.1842
```

```
In [17]: #Visualize Clusters for K = 3
kmeans = KMeans(n_clusters=3, random_state=42, n_init=10)
labels = kmeans.fit_predict(X_scaled)

plt.figure(figsize=(8,6))
sns.scatterplot(
    x=df_numeric['SALES'],
    y=df_numeric['MSRP'],
    hue=labels,
    palette='Set2'
)
plt.title("K-Means Clustering Visualization (K = 3)")
plt.xlabel("Sales")
plt.ylabel("MSRP")
plt.legend(title='Cluster')
plt.show()
```



In []: